

Prague University of Economics and Business
Faculty of Informatics and Statistics



A Data Integration Framework for Enterprise Application Security

MASTER THESIS

Study program: Applied Data Analytics and Artificial Intelligence

Specialization: Retail and E-Commerce Data Analytics

Author: Bc. Vojtěch Kraus

Supervisor: Ing. Aleš Kotuč

Prague, month YYYY

Acknowledgements

My sincere thanks go to Ing. Aleš Kotuč for his guidance and valuable feedback.

Abstract

To protect their business-critical applications from cybersecurity threats, enterprise security teams rely on numerous tools to detect and resolve vulnerabilities in application source code, configuration, dependencies, CI/CD pipelines, and runtime infrastructure. These tools use different integration patterns, protocols, and data formats, making it difficult to consistently parse, deduplicate, triage, and group their findings. In large environments, additional challenges emerge when linking all findings to the corresponding business applications. This thesis presents a data framework to tackle those challenges, delivering three distinct contributions: (1) a requirement specification for an application security data platform, (2) a reusable and extensible design covering architecture, data model, and medallion-based data pipeline, and (3) a reference implementation built on the Databricks platform. Together, these contributions offer application security teams a foundation for building a robust data platform independent of individual cybersecurity vendors, relieving them of risks usually associated with vendor lock-in.

Keywords

application security, software security, DevSecOps, security integration, data consolidation, medallion architecture, Databricks

Contents

Introduction	11
Motivation	11
Objective	11
Methodology	11
Structure	12
1 Literature Review	13
1.1 Software Development Literature	13
1.1.1 Source Code Management	13
1.1.2 Software Development Lifecycle	13
1.2 Application Security Literature	13
1.2.1 Code Security	13
1.2.2 Software Supply Chain Security	13
1.2.3 Software Lifecycle Security	13
1.2.4 Application Penetration Testing	13
1.2.5 Application Security Data Standards	14
1.3 Data Modeling and Engineering Literature	14
1.3.1 Data Architecture	14
1.3.2 Data Modeling and Pipelines	14
1.4 Related Work and Gap Analysis	14
2 Requirement Analysis and Specification	16
2.1 Personas	16
2.1.1 Application Owners	16
2.1.2 Application Developers	16
2.1.3 Security Experts	16
2.1.4 Development and Security Leadership	16
2.2 Data Sources	16
2.2.1 Tools for Software Development and Delivery	16
2.2.2 Runtime Infrastructure Platforms and Security Tools	17
2.2.3 Application Inventory	17
2.2.4 Security Testing Tools	17
2.3 Data Transformations	18
2.3.1 Data Ingestion and Validation	18
2.3.2 Data Normalization	18
2.3.3 Finding Deduplication	18
2.3.4 Business Application Mapping	18
2.3.5 Triage	18
2.4 Data Consumers	18

2.4.1	Remediation Tracking Systems	18
2.4.2	Application Security Posture Management (ASPM)	18
2.4.3	SOAR and SIEM	19
2.5	Non-Functional Requirements	19
2.5.1	Data Quality	19
2.5.2	Extensibility	19
2.5.3	Scalability	19
2.5.4	Performance	19
2.5.5	Security	19
2.5.6	Maintainability	19
2.6	Requirement Prioritization	19
3	Architecture	20
3.1	Architecture Context	20
3.2	Architecture Principles and Constraints	20
3.3	Technology Selection	20
3.4	High-Level Solution Architecture	21
3.5	Component Design	21
3.5.1	Data Ingestion Tier	21
3.5.2	Data Processing Tier	21
3.5.3	Data Serving Tier	22
3.6	Integration Architecture	22
3.6.1	Connector Strategy	22
3.7	Extensibility Architecture	23
3.8	Security Architecture	23
4	Pipeline and Model Design	24
4.1	Pipeline Overview	24
4.2	Bronze Layer	24
4.2.1	Transformations	24
4.2.2	Schema	25
4.3	Silver Layer	25
4.3.1	Transformations	25
4.3.2	Schema	25
4.4	Gold Layer	26
4.4.1	Transformations	26
4.4.2	Schema	26
5	Implementation	27
5.1	Environment Setup	27
5.2	Connector Implementation	27
5.3	Pipeline Implementation	27
5.4	Pipeline Orchestration	27
5.5	Data Serving	27

5.6	Testing and Validation	27
	Conclusion	28
	Thesis Outcomes and Contributions	28
	Limitations	28
	Future Work	28
	References	29
A	Generative AI Use Disclosure	31
A.1	Text Content	31
A.2	Diagrams and Figures	31
A.3	Source Code	32
A.4	Requirements Specification	32

List of Figures

- 3.1 Pipeline Architecture 22
- 4.1 Medallion Layer Design 24

List of Tables

List of source codes

3.1 [GitHub Security Connector Implementation](#) 20

Introduction

Motivation

Why I chose this topic:

- my appsec experience, what I see as a problem + back it up with some research
- explain that fundamentally, ASPM+VMDR are data problems, and having control over the data layer makes sense
- security vendor agnostic solution doesn't exist (there's DefectDojo but it's "just" a Django web app without proper data architecture, limited scaling abilities, no data pull capability, need to push with custom integration or upload results manually - DefectDojo Pro required for that)

Objective

Main goal: design and implement the model and pipeline (needs to be precisely formulated)

Subgoals:

- Summarize existing literature in ch1
- Analyze appsec domain, expectations from the framework, produce a requirements specification in ch2 and include full spec in App A
- In ch3, propose a reusable and extensible framework, domain driven design (name things as they are in appsec world), accommodating both OLTP and OLAP requirements, streaming for time critical use cases?
- Validate output of ch3 build a reference implementation of the framework in Databricks using Claude Code and Databricks AI Assistant
- Evaluate whether feasible/usable as a product + possible future improvements

Methodology

Methodology + outline hypotheses?

- First a quick review of existing literature and resources - include vendor documentation as it will need to be referenced later
- Requirement analysis - document req's for an appsec integration framework, what it should have

- Design, implement, evaluate the framework

Outline main results/outcomes of the thesis: business impact is that security teams can self-service deploy their data and ML pipeline and not rely on vendors' proprietary data design

Structure

1. Literature Review

1.1 Software Development Literature

What process/activities are we securing with appsec?

Books like Code Complete, Clean Code, Pragmatic Programmer...

1.1.1 Source Code Management

git book, software engineering at Google, ...

1.1.2 Software Development Lifecycle

Traditional vs devops and CI/CD pipelines, list books on each

1.2 Application Security Literature

In this section we focus on literature which explains the domain in which we're solving identified problems.

Books on general appsec

1.2.1 Code Security

1.2.2 Software Supply Chain Security

1.2.3 Software Lifecycle Security

1.2.4 Application Penetration Testing

Good source of wisdom for defensive teams - let's defend against what the attackers are doing

1.2.5 Application Security Data Standards

Standard data formats for vulnerabilities like Common Vulnerabilities and Exposures (CVE)

SARIF, CycloneDX, OCSF

Citing test (Anderson, [2020](#)), (Strengholt, [2025](#)),

Application Security Posture Management (ASPM) mentioned for the first time, ASPM mentioned second time.

Application Programming Interface (API)

1.3 Data Modeling and Engineering Literature

What tools do we have to achieve thesis goals?

1.3.1 Data Architecture

DWH

Lake

Lakehouse etc

1.3.2 Data Modeling and Pipelines

Data Modeling with Snowflake

DAMA-DMBOK

Medallion Architecture

1.4 Related Work and Gap Analysis

Security data integration - who wrote about it

An industry cybersecurity data model exists but only for SOC events, not for appsec modeling

Does DefectDojo have any useful models?

ASPM exists but not data oriented and not very customizable

2. Requirement Analysis and Specification

The goal of this chapter is to produce a requirements specification which we will later use to...

Appsec vendors, products and their capabilities, APIs and their maturity/completeness

Capture data formats and integration patterns of the tools

2.1 Personas

2.1.1 Application Owners

2.1.2 Application Developers

2.1.3 Security Experts

Architects and Engineers assigned to application teams

Including security compliance / auditors

2.1.4 Development and Security Leadership

Dashboards, metrics, trends, quantified risk

2.2 Data Sources

2.2.1 Tools for Software Development and Delivery

important data sources for inventory

start with SCM and issue tracking tools

Github, Gitlab, Bitbucket, Azure DevOps

Jira, Azure DevOps, Github and Gitlab issues

CI/CD pipelines - Jenkins, Github Actions, Azure DevOps

2.2.2 Runtime Infrastructure Platforms and Security Tools

We need infra portrayed from app angle - what does the app run on, how is it secure, whats the security context

AWS - EC2, serverless Azure GCP

Also how secure is it (infra VMDR like Qualys, Wiz...) can impact risk of app vulns if they run on vulnerable infra too

How to get data from Kubernetes, Docker even if on prem

2.2.3 Application Inventory

Start with looking at things from a business application perspective

Business application inventory + business impact assessment

source systems: Primary vulnerability discovery tools (secret scanning, , , ...)

2.2.4 Security Testing Tools

Secret Scanning Tools

Static Application Security Testing (SAST) Tools

Static Application Security Testing (SAST)

Software Composition Analysis (SCA) Tools

Software Composition Analysis (SCA)

Dynamic Application Security Testing (DAST) Tools

Dynamic Application Security Testing (DAST)

Penetration Testing

Need a way to upload manual findings, e.g. PDF report from external pentesters

2.3 Data Transformations

2.3.1 Data Ingestion and Validation

Filtering?

2.3.2 Data Normalization

Standard format for each object and finding type

2.3.3 Finding Deduplication

When 2 tools find the same thing, link it, don't create 2 issues

2.3.4 Business Application Mapping

2.3.5 Triage

2.4 Data Consumers

2.4.1 Remediation Tracking Systems

Jira, ServiceNow

2.4.2 Application Security Posture Management (ASPM)

Can pull from data platform already cleaned data instead of from each source separately

2.4.3 SOAR and SIEM

2.5 Non-Functional Requirements

2.5.1 Data Quality

Data validations and transformations on input are critical

2.5.2 Extensibility

Important feature of the whole framework

2.5.3 Scalability

it's intended for large enterprises, 10k's of repos +

2.5.4 Performance

Some use cases require low latency

2.5.5 Security

Obviously, with this use case

2.5.6 Maintainability

SaaS preferred over manual maintenance - Databricks is ideal

All source code versioned in git and following best practices

2.6 Requirement Prioritization

MoSCoW method

3. Architecture

generic concepts and open industry standards instead of proprietary concepts/technology unless necessary

Needs to support both OLTP/OLAP use cases based on requirements

3.1 Architecture Context

It will be used to pull data from data sources (those are naturally out of scope)

It needs to be ready to support an application security management dashboard with a combination of OLTP and OLAP use cases (dashboard application also out of scope).

3.2 Architecture Principles and Constraints

3.3 Technology Selection

What I'm choosing and why, what were the alternatives and reasons not to choose them

I will use Delta Live Tables (DLT), a data engineering pipeline framework running on top of Delta Lake that combines incremental ingestion, streamlined ETL, and automated data quality processes such as expectations (Lee et al., 2024). The advantage of DLT is...

Connectors - dlt, Fivetran?, Databricks Partners?

I use data load tool (dlt) to ingest data from security tools into the Bronze layer. The code snippet 3.1 shows an example implementation of a connector for GitHub security data.

Source code 3.1

GitHub Security Connector Implementation

```
1 import dlt
2 from dlt.sources.helpers.rest_client import paginate
3
4 @dlt.source
5 def github_security_source(org: str, token: str):
6     """Extract security data from GitHub organization."""
7
8     @dlt.resource(write_disposition="merge", primary_key="id")
9     def repos():
```

```

10     """Repository inventory"""
11     yield from paginate(f"/orgs/{org}/repos")
12
13     @dlt.resource(write_disposition="merge", primary_key="number")
14     def code_scanning_alerts():
15         """GitHub Code Scanning (SAST) alerts"""
16         yield from paginate(f"/orgs/{org}/code-scanning/alerts")
17
18     @dlt.resource(write_disposition="merge", primary_key="number")
19     def dependabot_alerts():
20         """Dependabot (SCA) vulnerability alerts"""
21         yield from paginate(f"/orgs/{org}/dependabot/alerts")
22
23     return repos, code_scanning_alerts, dependabot_alerts

```

Databricks

Unity Catalog

Delta Lake

Python vs Scala vs SQL

3.4 High-Level Solution Architecture

3.5 Component Design

3.5.1 Data Ingestion Tier

The various connectors

3.5.2 Data Processing Tier

Pipeline Architecture

Medallion architecture, see [Chapter 4](#)

DLT for defining transformations




Figure 3.1
Pipeline Architecture

Pipeline Orchestration

Databricks Workflows, don't think my workflows will be complex enough to use Airflow DAGs

3.5.3 Data Serving Tier

Lakebase

3.6 Integration Architecture

3.6.1 Connector Strategy

Different patterns and tools optimal for different data sources

As a service: Fivetran if it satisfies requirements (easiest to do, no need to reinvent the wheel)

Reuse existing: e.g. from DefectDojo <https://github.com/DefectDojo/django-DefectDojo/tree/master/doj>

Databricks Partner

dlt as default tool for custom connectors to REST APIs

3.7 Extensibility Architecture

Add new integrations via API

3.8 Security Architecture

4. Pipeline and Model Design

4.1 Pipeline Overview

Each medallion layer is implemented as a schema within a single Unity Catalog.




Figure 4.1
Medallion Layer Design

4.2 Bronze Layer

4.2.1 Transformations

Data source request filtering / configuration

Ingestion

4.2.2 Schema

by systems

4.3 Silver Layer

4.3.1 Transformations

Data Quality

Error handling

Normalization

Deduplication

4.3.2 Schema

Vendor-agnostic, unified for each entity and finding

Entities

Implemented as dimensions

Findings

Implemented as facts?

Relationships

(entity-to-entity, finding-to-entity)

4.4 Gold Layer

4.4.1 Transformations

ML Enrichment

Risk scoring

Triage

4.4.2 Schema

Aggregations

Metrics, counts, summaries...

Sliced by severity, app, owner, team, etc (requirement of grouping)

SLAs, trends, MTTR - what matters to stakeholders per stakeholder analysis

Read-Optimized Views

For dashboard queries, materialized, cached

5. Implementation

5.1 Environment Setup

Databricks on AWS

Workspace config

Users, groups, catalogs, schemas

5.2 Connector Implementation

Which connector which pattern and technology

What's the rate limiting on each? This is big data potentially

Authentication, pagination etc, don't forget to mention

5.3 Pipeline Implementation

5.4 Pipeline Orchestration

5.5 Data Serving

5.6 Testing and Validation

Conclusion

This thesis presented a data integration framework for enterprise application security...

Thesis Outcomes and Contributions

Evaluation against objectives from Introduction

Which requirements from Ch2 were met / partially met / not met

What did I deliver extra - ML models beyond requirements?

Limitations

What was not scoped / validated, constraints (license fees for commercial tools that don't give free dev trials)

What was difficult about implementing the design? What could not be implemented in the technology that I selected?

Future Work

More data sources / connectors, based on my provided blueprint

ML model improvements and fine-tuning

Potentials as an open source project

References

- Anderson, R. (2020). *Security engineering: A guide to building dependable distributed systems* (3rd ed.). John Wiley & Sons. (Cit. on p. [14](#)).
- Lee, D., Wentling, T., Haines, S., & Babu, P. (2024). *Delta Lake: The definitive guide: Modern data lakehouse architectures with data lakes*. O'Reilly Media. (Cit. on p. [20](#)).
- Strengtholt, P. (2025). *Building medallion architectures: Designing with Delta Lake and Spark*. O'Reilly Media. (Cit. on p. [14](#)).

Appendices

A. Generative AI Use Disclosure

This appendix contains disclosures on all usage of generative AI tools for this thesis.

A.1 Text Content

All textual content in this thesis was drafted and written exclusively by the author in Overleaf Online LaTeX editor with logged history of incremental manual edits. For language revision, Overleaf AI Assist tool with OpenAI GPT language model was used and customized with following prompt:

Article usage rules:

- Omit "the" before: plural nouns used generally, abstract concepts, technology names, proper nouns, section references
- Omit "a/an" when: introducing concepts already understood from context, in technical definitions, after "using/via/through"
- Keep articles only when: specific instance is referenced, first introduction of a countable noun, or omission creates ambiguity
- Examples of preferred style:
 - "Bronze layer stores raw data" (not "The Bronze layer stores the raw data")
 - "Figure 3 shows architecture" (not "Figure 3 shows the architecture")
 - "Data flows through pipeline" (not "The data flows through the pipeline")

A.2 Diagrams and Figures

All TikZ diagrams were generated by Claude Opus 4.5 (Anthropic) through iterative prompting. The author provided structural specifications, reviewed outputs, and requested modifications until diagrams accurately represented the intended architecture. An example prompt for Figure 3.1:

Create a TikZ pipeline architecture diagram with vertical flow:

Components:

1. Data Sources row: GitHub, SonarQube, Snyk, NVD, ...
2. External Fivetran box (outside Databricks boundary)
3. Ingestion tier: Fivetran Connector, dlt Connectors, Custom APIs
4. Processing tier: Bronze, Silver, Gold layers (colored)
5. Serving tier: Lakebase, REST API

6. Data Consumers row: Jira, ASPM, SIEM

Wrap Ingestion/Processing/Serving in Databricks Platform boundary.

Use arrows showing data flow between tiers.

Style: Soft colors, rounded corners, compact layout.

A.3 Source Code

Implementation code was written by the author with assistance from:

- **Claude Code** (Anthropic) with Opus 4.5 model for code generation, review, and debugging. Custom skills installed: vkraus/superpowers fork with embedded BDD skill and Databricks development skill.
- **Databricks AI Assistant** for in-editor code debugging and SQL optimization.

All generated code was reviewed, tested, and modified by the author before inclusion.

A.4 Requirements Specification

Requirements in Chapter ?? follow a test-centric traceability approach. Tests reference requirement IDs via pytest markers, and a generator script produces the traceability matrix by scanning test results. Implementation status is derived from test execution rather than manual tracking.